

Data — deep dive

Agenda

1. **Why this layer** — knowledge vs behaviour (RAG vs fine-tune)
2. **The pipeline** — sources → chunk → embed → store → retrieve → answer → **evaluate**
(*Phases 1–7*)
3. **RAG on production** — the advanced pipeline · **LangChain** · RAG vs long context · the **RAG platform**
4. **Labs** — build simple + advanced RAG on a free Kaggle **T4**

Why this layer

A base model is **frozen at its knowledge cutoff** and has never seen your private data. This layer supplements it with fresh / private knowledge **at answer time — no retraining**. The headline technique is **RAG** (Retrieval-Augmented Generation).


- Need new **knowledge**? → **RAG** (this layer)
- Need new **behaviour**? → **fine-tune** (`02-models/`)

The line that holds the course together: **RAG for knowledge, fine-tune for behaviour**. They stack — fine-tune the tone, retrieve the facts.

How we got here · the evolution of retrieval


Era 1 · Keyword — matches *words*, not meaning. A query for *"my pod keeps restarting"* never finds the doc *"fixing CrashLoopBackOff"* — zero shared words.

Era 2 · Semantic — turn text into a **vector**, so *meaning becomes position*: similar ideas land nearby even with no shared words.



The simple pipeline

INDEX · offline, once



write



ANSWER · per query



search

top-k chunks



Answer + cite

same embedder as INDEX ↑

Phase 1 - Sources

The raw knowledge. Its **shape** decides how much pipeline work before it's usable.

Shape	Examples	Pipeline cost
Unstructured	PDFs, web, email, transcripts	High — parse, clean, chunk
Semi-structured	Markdown, HTML, JSON, CSV	Medium — structure-aware split
Structured	SQL, APIs, spreadsheets	Low — often query directly

NIST SP 800-190

APPLICATION CONTAINER SECURITY GUIDE

- Recover: Recovery Planning
 - RC.RP-1: Recovery plan is executed during or after an event

Table 3 lists the security controls from NIST SP 800-53 Revision 4 [29] that can be accomplished partially or completely by using container technologies. The rightmost column lists the sections of this document that map to each NIST SP 800-53 control.

NIST SP 800-53 Control	Container Technology Relevancy	Related Sections of This Document
CM-3, Configuration Change Control	Images can be used to help manage change control for apps.	2.1, 2.2, 2.3, 2.4, 4.1
SC-2, Application Partitioning	Separating user functionality from administrator functionality can be accomplished in part by using containers or other virtualization technologies so that the functionality is performed in different containers.	2 (introduction), 2.3, 4.5.2
SC-3, Security Function Isolation	Separating security functions from non-security functions can be accomplished in part by using containers or other virtualization technologies so that the functions are performed in different containers.	2 (introduction), 2.3, 4.5.2
SC-4, Information in Shared Resources	Container technologies are designed to restrict each container's access to shared resources so that information cannot inadvertently be leaked from one container to another.	2 (introduction), 2.2, 2.3, 4.4
SC-6, Resource Availability	The maximum resources available for each container can be specified, thus protecting the availability of resources by not allowing any container to consume excessive resources.	2.2, 2.3
SC-7, Boundary Protection	Boundaries can be established and enforced between containers to restrict their communications with each other.	2 (introduction), 2.2, 2.3, 4.4
SC-39, Process Isolation	Multiple containers can run processes simultaneously on the same host, but those processes are isolated from each other.	2 (introduction), 2.1, 2.2, 2.3, 4.4
SI-7, Software, Firmware, and Information Integrity	Unauthorized changes to the contents of images can easily be detected and the altered image replaced with a known good copy.	2.3, 4.1, 4.2
SI-14, Non-Persistence	Images running within containers are replaced as needed with new image versions, so data, files, executables, and other information stored within running images is not persistent.	2.1, 2.3, 4.1

Similar to Table 3, Table 4 lists the NIST Cybersecurity Framework [30] subcategories that can be accomplished partially or completely by using container technologies. The rightmost column lists the sections of this document that map to each Cybersecurity Framework subcategory.

Cybersecurity Framework Subcategory	Container Technology Relevancy	Related Sections of This Document
-------------------------------------	--------------------------------	-----------------------------------

TLP: CLEAR

Figure 1: Incident Response Process

Preparation Phase

Prepare for major incidents *before they occur* to mitigate any impact on the organization. Preparation activities include:

- Documenting and understanding policies and procedures for incident response
- Instrumenting the environment to detect suspicious and malicious activity
- Establishing staffing plans
- Educating users on cyber threats and notification procedures
- Leveraging cyber threat intelligence (CTI) to proactively identify potential malicious activity

Define baseline systems and networks before an incident occurs to understand the basics of "normal" activity. Establishing baselines enables defenders to identify deviations. Preparation also includes

- Having infrastructure in place to handle complex incidents, including classified and out-of-band communications
- Developing and testing courses of action (COAs) for containment and eradication
- Establishing means for collecting digital forensics and other data or evidence

The goal of these items is to ensure resilient architectures and systems to maintain critical operations in a compromised state. Active defense measures that employ methods such as redirection and monitoring of adversary activities may also play a role in developing a robust incident response.

6

Phase 2 · clean & chunk

Clean (strip boilerplate), then **chunk** — passages small enough to embed, large enough to carry meaning. **The highest-leverage, most-overlooked choice in the layer.**

Strategy	When
Fixed-size (N tokens + overlap)	default, robust
Recursive (paragraph→sentence→word)	prose
Structure-aware (headings, code, rows)	docs / code
Semantic (split where topic shifts)	best quality, more compute

Two knobs: **chunk size** (tokens, 256–1024) and **overlap** (~10–20%). Store each chunk with **metadata** (source · page · date · ACL) → filter + cite.

Phase 3 · Embeddings

Text → **vector** that encodes *meaning* — near in meaning, near in space:

```
"how do I reset my password" → [ 0.21, -0.07, 0.88, ... ]  
"forgot my login credentials" → [ 0.19, -0.05, 0.85, ... ]  
cosine ≈ 0.97 → near in meaning → retrieve
```

Choose on **dimensions** · **max input length** · **domain/language** · **open vs API**.

Non-negotiable rule: embed documents *and* queries with the **same model**.

Change the model ⇒ re-embed the entire corpus.

Phase 4 · Vector store

A DB for millions of embeddings, answering *nearest-neighbour* in ms. Real stores use an **ANN** index — trade a sliver of recall for a huge speed-up.


Index	Idea	Trade-off
Flat	compare against every vector	exact, slow past ~100k
HNSW	navigable neighbour graph	fast + high recall; more memory (<i>default</i>)
IVF	cluster, search nearest clusters	smaller memory; tune <code>nprobe</code>

FAISS (*our labs*) · Chroma · pgvector · Qdrant/Weaviate/Milvus (production scale).

Metric must match the embedder (**cosine**/dot). Filter metadata **during** search.

Phase 5 · Retrieval

At answer time the query is embedded with the **same model**; the store returns the **top-k** nearest chunks — and *only* those become the context.



top-k (3–10): too few starves the model; too many buries the answer in noise + cost.

Phase 6 · RAG — putting it together

Stitch the retrieved chunks into the prompt so the model answers **from them**, with a citation back to the source:

```
SYSTEM: Answer only from the CONTEXT. If it isn't there, say you don't know.  
CONTEXT: <chunk 1> <chunk 2> ... <chunk k>  
USER:    <the question>
```

Phase 7 · Evaluation — measure **both** stages

"The demo answered well" is **not evidence**. RAG fails in two independent places, so you must measure — and debug — each separately:

Stage	Question it answers
① Retrieval	did we fetch the right chunks?
② Generation	did the answer use them faithfully?

Diagnose by stage — first ask: *was the correct chunk in the context?*

-  correct chunk retrieved **but the answer is wrong** → **generation** problem (prompt, model, or context too noisy)
-  correct chunk **was never retrieved** → **retrieval** problem (chunking, embedding, top-k, add a reranker)

Phase 7 · Retrieval metrics — *did we fetch the right chunks?*

Scored against a **gold set** (each question tagged with the chunk/doc that is truly relevant):

Metric	What it asks	Intuition
recall@k	of <i>all</i> relevant chunks, how many made the top-k?	did we miss anything?
precision@k	of the <i>k</i> chunks fetched, how many are relevant?	how much noise did we pull in?

Example: top-k = 5, and 2 of the 3 relevant chunks land in it → **recall@5 = 2/3 ≈ 0.67**;

if only those 2 of the 5 are relevant → **precision@5 = 2/5 = 0.4**.

Lab 2 measures retrieval on 40 hand-written questions mapped to a gold doc.

Phase 7 · Generation metrics & hallucination

Hallucination = a fluent, confident answer **not supported by the retrieved context** — the killer failure, because it looks *identical* to a correct answer.


Metric	What it asks
faithfulness	is every claim in the answer grounded in a chunk? (<i>this is what catches hallucination</i>)
answer relevance	does the answer actually address the question asked?

RAGAS scores these automatically with an **LLM-as-judge** — a second LLM reads *question + context + answer* and grades it, so you evaluate hundreds of cases without hand-reading each. Report quality **with latency + cost**, never a single number.

RAG on production

Beyond the naïve pipeline — better retrieval, a framework, and operating it at scale.

Advanced pipeline



Four levers (amber) — each closes a gap Chapter 1 left, measured against Phase 7:

- ① **Query understanding** — rewrite / multi-query / HyDE, don't trust the phrasing.
- ② **Hybrid** — dense vectors (meaning) ∪ **BM25** (exact terms: codes, flags, names).
- ③ **Rerank** — over-fetch 30, a cross-encoder keeps the best 5. *Biggest lever after chunking.*
- ④ **Metadata filter** — restrict by source · ACL · date.

LangChain

A general framework for LLM apps (agents, chatbots, RAG). It adds **no new concept** — it *productizes the seven phases*, one swappable component each.

Phase	By hand	LangChain
1 Sources	<code>{id,text,meta}</code> dicts	<code>Document</code>
2 Chunk	hand-written splitter	<code>RecursiveCharacterTextSplitter</code>
3 Embeddings	<code>SentenceTransformer.encode</code>	<code>HuggingFaceEmbeddings</code>
4 Vector store	<code>faiss.IndexFlatIP</code>	<code>FAISS.from_documents</code>
5 Retrieval	hand-written <code>retrieve()</code>	<code>store.as_retriever()</code>
6 RAG prompt	manual template + <code>.generate</code>	<code>ChatPromptTemplate</code> <code>chat</code>


Client-side orchestration; **eval + ops live outside the chain**. Proven in `labs/lab1-langchain/`.

RAG vs long context — why not just stuff the window?

Dimension	Long context	RAG
Infrastructure	none — "no-stack stack"	heavy: chunk + embed + vector DB + sync
Reliability	model sees everything	probabilistic → silent failure
Cost / query	reprocesses every token, every call	pays once at index time
Data ceiling	~1M tokens	effectively infinite corpus

- **Bounded data + global reasoning** (one contract, one book) → **long context**.
- **Fresh / private / huge** (an enterprise corpus) → **RAG**.

RAG platform



You don't write RAG — you operate a platform; each app is config on top.

Data-ops at fleet scale: freshness · cost · governance (ACL, right-to-delete).

Labs (free Kaggle T4)

Lab	Chapter	What you do
1 · Simple RAG, made visible	1	print the chunks a query loads, top-k, chunk size → grounded + cited Qwen answer
1' · LangChain build	3	Lab 1 rebuilt component-for-component, same numbers out — the proof
2 · Advanced RAG, real docs	2	clean + chunk real k8s docs, metadata + hybrid + reranker, measure hit@3 / MRR

Same constraints as the Models labs: pick **T4**, fp16 LLM, embedder + reranker on CPU, Internet **On**. Setup in `labs/README.md`.

Takeaways

Retrieval is the heart of RAG — most failures are retrieval failures.

RAG for knowledge, fine-tune for behaviour. Chunking + reranking are the biggest levers.

Measure both stages. A framework writes the chain; **you operate the platform.**